

Code ▾

# Machine Learning R Notebook

This is an R Markdown (<http://rmarkdown.rstudio.com>) Notebook. When you execute code within the notebook, the results appear beneath the code.

Hide

```
# check current working directory
getwd()
```

```
[1] "C:/Users/ajaka/OneDrive - Teesside University/CIS4035 - Machine Learning/E4216209_AJAKAYE_JESUBUKADE"
```

Hide

```
# list the files in the working directory
list.files(getwd())
```

```
[1] "caret_knn_base.rds"           "caret_knn_tune.rds"
[3] "caret_lr_base.rds"           "caret_lr_tune.rds"
[5] "caret_nb_base.rds"           "caret_nb_tune.rds"
[7] "caret_rf_base.rds"           "caret_rf_tune.rds"
[9] "E4216209_AJAKAYE_JESUBUKADE.nb.html" "E4216209_AJAKAYE_JESUBUKADE.Rmd"
[11] "Machine Learning R Notebook.pdf" "US_Accidents_March23_sampled_500k.csv"
```

## Importing Required Libraries and Reading Dataset

Hide

```
# Loading library
library(tidyverse) # for data processing
```

```
— Attaching core tidyverse packages ————— tidyverse 2.0.0 —
✓ dplyr     1.1.4    ✓ readr     2.1.5
✓ forcats   1.0.0    ✓ stringr   1.5.1
✓ ggplot2   3.5.1    ✓ tibble    3.2.1
✓ lubridate 1.9.4    ✓ tidyr     1.3.1
✓ purrr    1.0.2    — Conflicts ————— tidyvers
e_conflicts() —
✖ dplyr::filter() masks stats::filter()
✖ dplyr::lag()    masks stats::lag()
ℹ Use the conflict_prefer("tidyverse") to force all conflicts to become errors
```

Hide

```
library(gridExtra) # merging visuals
```

```
Attaching package: 'gridExtra'
```

```
The following object is masked from 'package:dplyr':
```

```
combine
```

[Hide](#)

```
library(ROSE) # Data Balancing
```

```
Loaded ROSE 0.0-4
```

[Hide](#)

```
library(caret) # Model Building (createDataPartition)
```

```
Loading required package: lattice
```

```
Attaching package: 'caret'
```

```
The following object is masked from 'package:purrr':
```

```
lift
```

[Hide](#)

```
library(pROC) # ROC and AUC
```

```
Type 'citation("pROC")' for a citation.
```

```
Attaching package: 'pROC'
```

```
The following objects are masked from 'package:stats':
```

```
cov, smooth, var
```

[Hide](#)

```
library(MLmetrics) # F1 score
```

```
Warning: package 'MLmetrics' was built under R version 4.4.3
```

```
Attaching package: 'MLmetrics'
```

```
The following objects are masked from 'package:caret':
```

```
MAE, RMSE
```

```
The following object is masked from 'package:base':
```

```
Recall
```

[Hide](#)

```
# Read the csv file
US_data <- read.csv("US_Accidents_March23_sampled_500k.csv", na.strings=c("", "NA"))
head(US_data)
```

ID	Source	Severity	Start_Time	End_Time
<chr>	<chr>	<int>	<chr>	<chr>
1 A-2047758	Source2	2	2019-06-12 10:10:56	2019-06-12 10:55:58
2 A-4694324	Source1	2	2022-12-03 23:37:14.0000000000	2022-12-04 01:56:53.00000000
3 A-5006183	Source1	2	2022-08-20 13:13:00.0000000000	2022-08-20 15:22:45.00000000
4 A-4237356	Source1	2	2022-02-21 17:43:04	2022-02-21 19:43:23
5 A-6690583	Source1	2	2020-12-04 01:46:00	2020-12-04 04:13:09
6 A-1101469	Source2	2	2021-03-29 07:03:58	2021-03-29 08:51:01

6 rows | 1-6 of 46 columns

## Data Preparation and Data Cleaning

[Hide](#)

```
# Get the column names of data
colnames(US_data)
```

```
[1] "ID"                      "Source"                  "Severity"
[4] "Start_Time"               "End_Time"                "Start_Lat"
[7] "Start_Lng"                "End_Lat"                 "End_Lng"
[10] "Distance.mi."             "Description"              "Street"
[13] "City"                     "County"                  "State"
[16] "Zipcode"                  "Country"                 "Timezone"
[19] "Airport_Code"              "Weather_Timestamp"       "Temperature.F."
[22] "Wind_Chill.F."            "Humidity..."              "Pressure.in."
[25] "Visibility.mi."            "Wind_Direction"          "Wind_Speed.mph."
[28] "Precipitation.in."         "Weather_Condition"        "Amenity"
[31] "Bump"                     "Crossing"                "Give_Way"
[34] "Junction"                 "No_Exit"                 "Railway"
[37] "Roundabout"                "Station"                 "Stop"
[40] "Traffic_Calming"           "Traffic_Signal"           "Turning_Loop"
[43] "Sunrise_Sunset"              "Civil_Twilight"           "Nautical_Twilight"
[46] "Astronomical_Twilight"
```

[Hide](#)

```
# Checking Data attributes
str(US_data)
```

```
'data.frame': 500000 obs. of 46 variables:
 $ ID           : chr "A-2047758" "A-4694324" "A-5006183" "A-4237356" ...
 $ Source       : chr "Source2" "Source1" "Source1" "Source1" ...
 $ Severity     : int 2 2 2 2 2 2 2 2 2 2 ...
 $ Start_Time   : chr "2019-06-12 10:10:56" "2022-12-03 23:37:14.000000000" "2022-08
-20 13:13:00.000000000" "2022-02-21 17:43:04" ...
 $ End_Time     : chr "2019-06-12 10:55:58" "2022-12-04 01:56:53.000000000" "2022-08
-20 15:22:45.000000000" "2022-02-21 19:43:23" ...
 $ Start_Lat    : num 30.6 39 34.7 43.7 35.4 ...
 $ Start_Lng    : num -91.2 -77.4 -120.5 -93 -119 ...
 $ End_Lat      : num NA 39 34.7 43.7 35.4 ...
 $ End_Lng      : num NA -77.4 -120.5 -93 -119 ...
 $ Distance.mi. : num 0 0.056 0.022 1.054 0.046 ...
 $ Description   : chr "Accident on LA-19 Baker-Zachary Hwy at Lower Zachary Rd." "In
cident on FOREST RIDGE DR near PEPPERIDGE PL Drive with caution." "Accident on W Central Ave
from Floradale Ave to Western Ave." "Incident on I-90 EB near REST AREA Drive with caution."
...
$ Street        : chr "Highway 19" "Forest Ridge Dr" "Floradale Ave" "14th St NW"
...
$ City          : chr "Zachary" "Sterling" "Lompoc" "Austin" ...
$ County        : chr "East Baton Rouge" "Loudoun" "Santa Barbara" "Mower" ...
$ State          : chr "LA" "VA" "CA" "MN" ...
$ Zipcode        : chr "70791-4610" "20164-2813" "93436" "55912" ...
$ Country        : chr "US" "US" "US" "US" ...
$ Timezone       : chr "US/Central" "US/Eastern" "US/Pacific" "US/Central" ...
$ Airport_Code   : chr "KBTR" "KIAD" "KLPC" "KAUM" ...
$ Weather_Timestamp : chr "2019-06-12 09:53:00" "2022-12-03 23:52:00" "2022-08-20 12:56:
00" "2022-02-21 17:35:00" ...
$ Temperature.F. : num 77 45 68 27 42 42 35 90 91 63 ...
$ Wind_Chill.F.  : num 77 43 68 15 42 35 35 90 91 63 ...
$ Humidity...     : num 62 48 73 86 34 58 89 55 39 78 ...
$ Pressure.in.   : num 29.9 29.9 29.8 28.5 29.8 ...
$ Visibility.mi. : num 10 10 10 10 10 10 10 10 10 10 ...
$ Wind_Direction : chr "NW" "W" "W" "ENE" ...
$ Wind_Speed.mph. : num 5 5 13 15 0 13 0 12 7 10 ...
$ Precipitation.in. : num 0 0 0 0 0 0 0 0 0 0 ...
$ Weather_Condition : chr "Fair" "Fair" "Fair" "Wintry Mix" ...
$ Amenity        : chr "False" "False" "False" "False" ...
$ Bump           : chr "False" "False" "False" "False" ...
$ Crossing        : chr "False" "False" "False" "False" ...
$ Give_Way        : chr "False" "False" "False" "False" ...
$ Junction        : chr "False" "False" "False" "False" ...
$ No_Exit         : chr "False" "False" "False" "False" ...
$ Railway          : chr "False" "False" "False" "False" ...
$ Roundabout       : chr "False" "False" "False" "False" ...
$ Station          : chr "False" "False" "False" "False" ...
$ Stop             : chr "False" "False" "False" "False" ...
$ Traffic_Calming : chr "False" "False" "False" "False" ...
$ Traffic_Signal   : chr "True" "False" "True" "False" ...
$ Turning_Loop     : chr "False" "False" "False" "False" ...
$ Sunrise_Sunset   : chr "Day" "Night" "Day" "Day" ...
$ Civil_Twilight   : chr "Day" "Night" "Day" "Day" ...
$ Nautical_Twilight: chr "Day" "Night" "Day" "Day" ...
$ Astronomical_Twilight: chr "Day" "Night" "Day" "Day" ...
```

[Hide](#)

```
# Checking for missing values
sum(is.na(US_data))
```

[1] 829873

[Hide](#)

```
# Checking for proportion of missing data
US_pctmiss <- colSums(is.na(US_data))/nrow(US_data)
round(US_pctmiss, 2)
```

ID	Source	Severity
0.00	0.00	0.00
Start_Time	End_Time	Start_Lat
0.00	0.00	0.00
Start_Lng	End_Lat	End_Lng
0.00	0.44	0.44
Distance.mi.	Description	Street
0.00	0.00	0.00
City	County	State
0.00	0.00	0.00
Zipcode	Country	Timezone
0.00	0.00	0.00
Airport_Code	Weather_Timestamp	Temperature.F.
0.00	0.02	0.02
Wind_Chill.F.	Humidity...	Pressure.in.
0.26	0.02	0.02
Visibility.mi.	Wind_Direction	Wind_Speed.mph.
0.02	0.02	0.07
Precipitation.in.	Weather_Condition	Amenity
0.29	0.02	0.00
Bump	Crossing	Give_Way
0.00	0.00	0.00
Junction	No_Exit	Railway
0.00	0.00	0.00
Roundabout	Station	Stop
0.00	0.00	0.00
Traffic_Calming	Traffic_Signal	Turning_Loop
0.00	0.00	0.00
Sunrise_Sunset	Civil_Twilight	Nautical_Twilight
0.00	0.00	0.00
Astronomical_Twilight		
0.00		

[Hide](#)

```
# Dropping All Rows with Missing Data
US_newdata <- na.omit(US_data)

# Re-coding Target variable
US_newdata <- mutate(US_newdata, Severity = ifelse(Severity < 3, 0, 1))


table(US_newdata$Severity)
```

0	1
218279	11648

[Hide](#)

```
# Add a plot of target here
target <- ggplot(US_newdata, aes(x = factor(Severity),
                                    fill = factor(Severity))) +
  geom_bar() +
  labs(title = "Bar Plot of Target Variable",
       x = "Severity",
       y = "Count",
       fill = "Severity") +
  scale_x_discrete(labels = c("0" = "LOW", "1" = "HIGH")) +
  scale_fill_discrete(labels = c("0" = "LOW", "1" = "HIGH")) +
  theme_minimal()

target
```



## Dealing with imbalance

[Hide](#)

```
# Undersampling
set.seed(199)
US_newdata_Sample = ovun.sample(Severity ~ .,
                                 data = US_newdata,
                                 method = "under",
                                 N = 25000)$data

# Checking for missing values
anyNA(US_newdata_Sample)
```

```
[1] FALSE
```

[Hide](#)

```
table(US_newdata_Sample$Roundabout)
```

```
False
25000
```

[Hide](#)

```
table(US_newdata_Sample$Turning_Loop)
```


```
False
25000
```

[Hide](#)


```
# Add a plot of new target here
# Add a plot of target here
balances_target <- ggplot(US_newdata_Sample, aes(x = factor(Severity),
                                                fill = factor(Severity))) +
  geom_bar() +
  labs(title = "Bar Plot of Target Variable",
       x = "Severity",
       y = "Count",
       fill = "Severity") +
  scale_x_discrete(labels = c("0" = "LOW", "1" = "HIGH")) +
  scale_fill_discrete(labels = c("0" = "LOW", "1" = "HIGH")) +
  theme_minimal()

# Plot Grid of Target Variable
grid.arrange(
  target,
  balances_target,
  ncol = 2
)
```

Bar Plot of Target Variable



Bar Plot of Target Variable



## Data Transformation

[Hide](#)

```
# Converting Date Columns to date data
US_newdata_Sample$Start_Time <- parse_date_time(US_newdata_Sample$Start_Time, orders = "ymd_HMS")

US_newdata_Sample$End_Time <- parse_date_time(US_newdata_Sample$End_Time, orders = "ymd_HMS")

US_newdata_Sample$Weather_Timestamp <- parse_date_time(US_newdata_Sample$Weather_Timestamp, orders = "ymd_HMS")

# Creating New Calculated Columns
US_newdata_Sample$Time <- (US_newdata_Sample$End_Time - US_newdata_Sample$Start_Time)

# Converting day difference to numeric
US_newdata_Sample$Time <- as.numeric(US_newdata_Sample$Time)
```

## Feature Selection 1

The following features are dropped for various reasons as outlined below:

1. ID, Source (Do not contain information pertaining to target variable)
2. Start\_Time, End\_Time (Transformed to new column "Time" showing total accident time)
3. Start\_Lat, Start\_Lng, End\_Lat, End\_Lng (Dropped as "Distance.mi" is in the dataset)
4. Country, Turning\_Loop, Roundabout (Only has one value)
5. Description (Unique values for each entry)

[Hide](#)

```
# Drop Columns not needed based on factors such as in already calculated columns
US_newdata_Sample <- select(US_newdata_Sample, -ID, -Source, -Start_Time, -End_Time, -Description, -Start_Lat, -Start_Lng, -End_Lat, -End_Lng, -Country, -Turning_Loop, -Roundabout)
```

## Data Transformation and Summaries

[Hide](#)

```
# Creating list of Columns
numeric_columns = names((which(sapply(select(US_newdata_Sample, -"Severity"),
                                         is.numeric)))))

numeric_columns
```

```
[1] "Distance.mi."      "Temperature.F."    "Wind_Chill.F."    "Humidity..." 
[5] "Pressure.in."     "Visibility.mi."   "Wind_Speed.mph."  "Precipitation.in."
[9] "Time"
```

[Hide](#)

```
categorical_columns <- names((which(sapply(US_newdata_Sample, is.character)))))

categorical_columns
```

```
[1] "Street"           "City"             "County"
[4] "State"            "Zipcode"          "Timezone"
[7] "Airport_Code"     "Wind_Direction"   "Weather_Condition"
[10] "Amenity"          "Bump"             "Crossing"
[13] "Give_Way"         "Junction"        "No_Exit"
[16] "Railway"          "Station"          "Stop"
[19] "Traffic_Calming" "Traffic_Signal"  "Sunrise_Sunset"
[22] "Civil_Twilight"   "Nautical_Twilight" "Astronomical_Twilight"
```

[Hide](#)


```
# check_categorical_columns <- names((which(sapply(US_newdata_Sample, is.character)))))

# check_categorical_columns
```

[Hide](#)

```
# create box plot for numeric variables
plots <- lapply(numeric_columns, function(colname) {
  ggplot(US_newdata_Sample, aes(y = .data[[colname]])) +
    geom_boxplot(fill = "lightgreen") +
    theme_minimal() +
    labs(title = colname, y = NULL) +
    coord_flip() # ← This flips the plot horizontally
})

# Arrange in 3x3 grid
do.call(grid.arrange, c(plots, ncol = 3))
```



```
# create side by side bar chart for categorical variables
Amenity <- ggplot(US_newdata_Sample, aes(x = Amenity,
                                         fill = factor(Severity))) +
  geom_bar(position = "fill") +
  labs(x = "Amenity",
       y = "Count",
       fill = "Severity") +
  scale_fill_discrete(labels = c("0" = "LOW", "1" = "HIGH")) +
  theme_minimal()

Bump <- ggplot(US_newdata_Sample, aes(x = Bump,
                                         fill = factor(Severity))) +
  geom_bar(position = "fill") +
  labs(x = "Bump",
       y = "Count",
       fill = "Severity") +
  scale_fill_discrete(labels = c("0" = "LOW", "1" = "HIGH")) +
  theme_minimal()

Crossing <- ggplot(US_newdata_Sample, aes(x = Crossing,
                                         fill = factor(Severity))) +
  geom_bar(position = "fill") +
  labs(x = "Crossing",
       y = "Count",
       fill = "Severity") +
  scale_fill_discrete(labels = c("0" = "LOW", "1" = "HIGH")) +
  theme_minimal()

Give_Way <- ggplot(US_newdata_Sample, aes(x = Give_Way,
                                         fill = factor(Severity))) +
  geom_bar(position = "fill") +
  labs(x = "Give_Way",
       y = "Count",
       fill = "Severity") +
  scale_fill_discrete(labels = c("0" = "LOW", "1" = "HIGH")) +
  theme_minimal()

Junction <- ggplot(US_newdata_Sample, aes(x = Junction,
                                         fill = factor(Severity))) +
  geom_bar(position = "fill") +
  labs(x = "Junction",
       y = "Count",
       fill = "Severity") +
  scale_fill_discrete(labels = c("0" = "LOW", "1" = "HIGH")) +
  theme_minimal()

No_Exit <- ggplot(US_newdata_Sample, aes(x = No_Exit,
                                         fill = factor(Severity))) +
  geom_bar(position = "fill") +
  labs(x = "No_Exit",
       y = "Count",
       fill = "Severity") +
  scale_fill_discrete(labels = c("0" = "LOW", "1" = "HIGH")) +
  theme_minimal()
```

```
Railway <- ggplot(US_newdata_Sample, aes(x = Railway,
                                             fill = factor(Severity))) +
  geom_bar(position = "fill") +
  labs(x = "Railway",
       y = "Count",
       fill = "Severity") +
  scale_fill_discrete(labels = c("0" = "LOW", "1" = "HIGH")) +
  theme_minimal()

Station <- ggplot(US_newdata_Sample, aes(x = Station,
                                             fill = factor(Severity))) +
  geom_bar(position = "fill") +
  labs(x = "Station",
       y = "Count",
       fill = "Severity") +
  scale_fill_discrete(labels = c("0" = "LOW", "1" = "HIGH")) +
  theme_minimal()


Stop <- ggplot(US_newdata_Sample, aes(x = Stop,
                                         fill = factor(Severity))) +
  geom_bar(position = "fill") +
  labs(x = "Stop",
       y = "Count",
       fill = "Severity") +
  scale_fill_discrete(labels = c("0" = "LOW", "1" = "HIGH")) +
  theme_minimal()

Traffic_Calming <- ggplot(US_newdata_Sample, aes(x = Traffic_Calming,
                                                 fill = factor(Severity))) +
  geom_bar(position = "fill") +
  labs(x = "Traffic_Calming",
       y = "Count",
       fill = "Severity") +
  scale_fill_discrete(labels = c("0" = "LOW", "1" = "HIGH")) +
  theme_minimal()

Traffic_Signal <- ggplot(US_newdata_Sample, aes(x = Traffic_Signal,
                                                 fill = factor(Severity))) +
  geom_bar(position = "fill") +
  labs(x = "Traffic_Signal",
       y = "Count",
       fill = "Severity") +
  scale_fill_discrete(labels = c("0" = "LOW", "1" = "HIGH")) +
  theme_minimal()

# Arrange in 3x3 grid
grid.arrange(
  Amenity,
  Bump,
  Crossing,
  Give_Way,
  Junction,
  No_Exit,
  Railway,
  Station,
```

```
Stop,
Traffic_Calming,
Traffic_Signal,
ncol = 3
)
```



Hide

```
# Encoding categorical Variables
US_newdata_Sample[categorical_columns] <- lapply(
  US_newdata_Sample[categorical_columns],
  function(x) as.integer(as.factor(x)) - 1
)

head(US_newdata_Sample)
```

	Severity	Distance.mi.	Street	City	Cou...	State	Zipcode	Timezone	Airport_Code
	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
1	0	0.019	6889	27	9	4	10395	2	89
2	0	1.007	7282	3614	605	2	10697	2	469
3	0	1.060	8894	3558	450	43	3316	1	1104
4	0	1.027	8960	147	364	9	5155	1	499
5	0	0.039	2067	935	261	41	9639	0	332
6	0	0.101	6521	2532	637	8	6119	1	847

6 rows | 1-10 of 35 columns

Hide

```
str(US_newdata_Sample)
```

```
'data.frame': 25000 obs. of 35 variables:
 $ Severity           : num  0 0 0 0 0 0 0 0 0 ...
 $ Distance.mi.       : num  0.019 1.007 1.06 1.027 0.039 ...
 $ Street              : num  6889 7282 8894 8960 2067 ...
 $ City                : num  27 3614 3558 147 935 ...
 $ County              : num  9 605 450 364 261 637 216 320 724 412 ...
 $ State               : num  4 2 43 9 41 8 9 3 8 38 ...
 $ Zipcode             : num  10395 10697 3316 5155 9639 ...
 $ Timezone            : num  2 2 1 1 0 1 1 3 1 1 ...
 $ Airport_Code         : num  89 469 1104 499 332 ...
 $ Weather_Timestamp   : POSIXct, format: "2023-01-20 22:01:00" "2021-09-16 19:51:00" ...
 $ Temperature.F.      : num  20 101 49 37 52 73 51 57 66 61 ...
 $ Wind_Chill.F.       : num  13 101 45 37 52 73 51 57 66 61 ...
 $ Humidity...          : num  89 15 26 73 69 46 21 63 45 30 ...
 $ Pressure.in.         : num  22.6 28.5 30 29.3 29.4 ...
 $ Visibility.mi.       : num  3 10 10 7 10 10 10 10 10 10 ...
 $ Wind_Direction       : num  3 8 19 0 8 10 10 1 7 11 ...
 $ Wind_Speed.mph.      : num  5 6 9 0 6 6 9 7 13 7 ...
 $ Precipitation.in.    : num  0.01 0 0 0 0 0 0 0 0 0 ...
 $ Weather_Condition   : num  33 7 43 7 7 39 39 7 7 7 ...
 $ Amenity              : num  0 0 0 0 0 0 0 0 0 0 ...
 $ Bump                 : num  0 0 0 0 0 0 0 0 0 0 ...
 $ Crossing             : num  1 0 0 0 0 0 0 0 0 1 ...
 $ Give_Way              : num  0 0 0 0 0 0 0 0 0 0 ...
 $ Junction              : num  0 0 0 0 1 0 1 0 0 0 ...
 $ No_Exit              : num  0 0 0 0 0 0 0 0 0 0 ...
 $ Railway              : num  0 0 0 0 0 0 0 0 0 0 ...
 $ Station               : num  0 0 0 0 0 0 0 0 0 0 ...
 $ Stop                 : num  0 0 0 0 0 0 0 0 0 0 ...
 $ Traffic_Calming      : num  0 0 0 0 0 0 0 0 0 0 ...
 $ Traffic_Signal        : num  1 0 0 0 0 0 0 0 0 0 ...
 $ Sunrise_Sunset         : num  1 1 0 1 1 0 0 1 0 1 ...
 $ Civil_Twilight         : num  1 1 0 1 1 0 0 1 0 1 ...
 $ Nautical_Twilight     : num  1 1 0 1 1 0 0 1 0 0 ...
 $ Astronomical_Twilight: num  1 1 0 1 1 0 0 1 0 0 ...
 $ Time                  : num  178.2 35.5 76.8 149.9 120 ...
```

Hide

```
# Standardizing Feature Columns
US_newdata_Sample <- US_newdata_Sample %>% mutate(Weather_Timestamp =
scale(Weather_Timestamp))

# Scale Categorical Columns
US_newdata_Sample[categorical_columns] <-
lapply(US_newdata_Sample[categorical_columns], scale)

# Scale Numeric Columns
normalise <-function(x) { (x -min(x))/(max(x)-min(x)) }

US_newdata_Sample[numerical_columns] <- as.data.frame(
lapply(US_newdata_Sample[,numerical_columns], normalise))

head(US_newdata_Sample)
```

	<b>Severity</b> <dbl>	<b>Distance.mi.</b> <dbl>	<b>Street</b> <dbl>	<b>City</b> <dbl>	<b>County</b> <dbl>	<b>State</b> <dbl>	<b>Zipcode</b> <dbl>	T
1	0	0.0001269290	0.02311716	-1.7319073	-1.9327793	-1.0913724	0.8598495	0.
2	0	0.0067272363	0.15227865	1.1072578	0.1318473	-1.2240307	0.9373198	0.
3	0	0.0070813014	0.68207081	1.0629329	-0.4050942	1.4954657	-0.9560843	-0.
4	0	0.0068608457	0.70376205	-1.6369255	-0.7030101	-0.7597265	-0.4843365	-0.
5	0	0.0002605384	-1.56165815	-1.0132115	-1.0598163	1.3628073	0.6659173	-1.
6	0	0.0006747278	-0.09782795	0.2508383	0.2426998	-0.8260557	-0.2370473	-0.

6 rows | 1-9 of 35 columns

## Feature Selection 2


### Correlation

[Hide](#)

```
corrmatrix <- cor(US_newdata_Sample)
# Generate new variable names
var_codes <- paste0("V", seq_along(US_newdata_Sample))
names(var_codes) <- US_newdata_Sample # This will be used for the key


# Rename rows and columns of the correlation matrix
colnames(corrmatrix) <- var_codes
rownames(corrmatrix) <- var_codes

corrplot::corrplot(corrmatrix, method = "color", type = "upper", tl.srt = 45)
```

[Hide](#)

```
ggplot(US_newdata_Sample, aes(x = Wind_Chill.F., y = Temperature.F.)) +  
  geom_point(color = "steelblue", alpha = 0.6, size = 2.5) +  
  geom_smooth(method = "lm", se = FALSE, color = "darkred") +  
  labs(title = "Scatter Plot with Trend Line",  
       x = "Wind_Chill.F",  
       y = "Temperature.F") +  
  theme_minimal()
```

## Scatter Plot with Trend Line


[Hide](#)

```
# correlation for all variables
correlations = round(cor(US_newdata_Sample), digits = 2)

# Get variable greater than 0.8
weights <- which(abs(correlations) > 0.8 & row(correlations)<col(correlations), arr.ind=TRUE)

## reconstruct names from positions
high_cor <- matrix(colnames(correlations)[weights], ncol=2)

high_cor
```

```
[,1]      [,2]
[1,] "Temperature.F."  "Wind_Chill.F."
[2,] "Sunrise_Sunset"   "Civil_Twilight"
[3,] "Sunrise_Sunset"   "Nautical_Twilight"
[4,] "Civil_Twilight"   "Nautical_Twilight"
[5,] "Civil_Twilight"   "Astronomical_Twilight"
[6,] "Nautical_Twilight" "Astronomical_Twilight"
```

[Hide](#)

```
##### Remove Based on correlation
US_newdata_Sample <- select(US_newdata_Sample, -Wind_Chill.F.,
                           -Nautical_Twilight, -Astronomical_Twilight,
                           -Civil_Twilight)
```

## MultiCollinearity using VIF

[Hide](#)

```
# Checking for MultiCollinearity using VIF
set.seed(123)
model_data_glm <- glm(Severity ~ ., data = US_newdata_Sample, family = binomial)
model_data_vif <- car::vif(model_data_glm)
model_data_vif
```

Distance.mi.	Street	City	County
1.042293	1.120201	1.116959	1.081503
State	Zipcode	Timezone	Airport_Code
1.212660	1.535404	1.477869	1.124924
Weather_Timestamp	Temperature.F.	Humidity...	Pressure.in.
1.129427	1.369699	1.557705	1.234861
Visibility.mi.	Wind_Direction	Wind_Speed.mph.	Precipitation.in.
1.357769	1.216594	1.317629	1.100350
Weather_Condition	Amenity	Bump	Crossing
1.099975	1.043610	1.997277	1.421154
Give_Way	Junction	No_Exit	Railway
1.013457	1.069210	1.005325	1.054065
Station	Stop	Traffic_Calming	Traffic_Signal
1.080898	1.036682	1.997021	1.343610
Sunrise_Sunset	Time		
1.251108	1.007910		

## Building Base Model and Testing

Hide

```
# Checking for missing values
anyNA(US_newdata_Sample)
```

```
[1] FALSE
```

Hide

```
# Checking target distribution
table(US_newdata_Sample$Severity)
```

```
0      1
13352 11648
```

Hide

```
# factor target
US_newdata_Sample$Severity <- as.factor(US_newdata_Sample$Severity)
```

Hide

```
# Splitting Data into Training and Testing Set
set.seed(42)
train_index <- createDataPartition(US_newdata_Sample$Severity, p = 0.8, list = FALSE)
train_data <- US_newdata_Sample[train_index, ]
test_data <- US_newdata_Sample[-train_index, ]
```

## KNN

[Hide](#)

```
set.seed(42)
start_time <- Sys.time() # Record start time for model
model_knn <- train(Severity ~ ., data = train_data, method = "knn")
stop_time <- Sys.time() # Record model ending time
training_time <- stop_time - start_time # Calculate training time

cat("Training Time: ", training_time, "\n")
```

Training Time: 1.187466

[Hide](#)

```
print(model_knn$bestTune)
```

Error: object 'model\_knn' not found

[Hide](#)

```
# Load KNN base model
model_knn <- readRDS("caret_knn_base.rds")
model_knn
```

k-Nearest Neighbors

20001 samples  
30 predictor  
2 classes: '0', '1'

No pre-processing  
Resampling: Bootstrapped (25 reps)  
Summary of sample sizes: 20001, 20001, 20001, 20001, 20001, 20001, ...  
Resampling results across tuning parameters:

k	Accuracy	Kappa
5	0.6923616	0.3784880
7	0.6994544	0.3922544
9	0.7029754	0.3989455

Accuracy was used to select the optimal model using the largest value.  
The final value used for the model was k = 9.

[Hide](#)

```
# Base KNN Test Prediction  
knn_test_predict <- predict(model_knn, test_data)  
confusionMatrix(knn_test_predict, test_data$Severity)
```

### Confusion Matrix and Statistics

		Reference
Prediction	0	1
0	2162	804
1	508	1525

Accuracy : 0.7375  
95% CI : (0.7251, 0.7497)  
No Information Rate : 0.5341  
P-Value [Acc > NIR] : < 2.2e-16  
  
Kappa : 0.4683  
  
McNemar's Test P-Value : 3.814e-16  
  
Sensitivity : 0.8097  
Specificity : 0.6548  
Pos Pred Value : 0.7289  
Neg Pred Value : 0.7501  
Prevalence : 0.5341  
Detection Rate : 0.4325  
Detection Prevalence : 0.5933  
Balanced Accuracy : 0.7323  
  
'Positive' Class : 0

## Naive Bayes

Hide

```
# Load NB base model  
model_nb <- readRDS("caret_nb_base.rds")  
model_nb
```

## Naive Bayes

```
20001 samples
 30 predictor
 2 classes: '0', '1'
```

No pre-processing

Resampling: Bootstrapped (25 reps)

Summary of sample sizes: 20001, 20001, 20001, 20001, 20001, 20001, ...

Resampling results across tuning parameters:

usekernel	Accuracy	Kappa
FALSE	0.6682345	0.3390938
TRUE	0.6884272	0.3808599

Tuning parameter 'fL' was held constant at a value of 0

Tuning parameter 'adjust'

was held constant at a value of 1

Accuracy was used to select the optimal model using the largest value.

The final values used for the model were fL = 0, usekernel = TRUE and adjust = 1.

[Hide](#)

```
# Base Naive Bayes Test Prediction
nb_test_predict <- predict(model_nb, test_data)
confusionMatrix(nb_test_predict, test_data$Severity)
```

## Confusion Matrix and Statistics

Reference	0	1
Prediction 0	2009	761
1	661	1568

Accuracy : 0.7155  
 95% CI : (0.7028, 0.728)

No Information Rate : 0.5341  
 P-Value [Acc > NIR] : < 2.2e-16

Kappa : 0.4269

McNemar's Test P-Value : 0.008656

Sensitivity : 0.7524  
 Specificity : 0.6733  
 Pos Pred Value : 0.7253  
 Neg Pred Value : 0.7035  
 Prevalence : 0.5341  
 Detection Rate : 0.4019  
 Detection Prevalence : 0.5541  
 Balanced Accuracy : 0.7128

'Positive' Class : 0

# Logistic Regression

[Hide](#)

```
# Load LR base model
model_lr <- readRDS("caret_lr_base.rds")
model_lr
```

Generalized Linear Model

```
20001 samples
30 predictor
2 classes: '0', '1'
```

No pre-processing

Resampling: Bootstrapped (25 reps)

Summary of sample sizes: 20001, 20001, 20001, 20001, 20001, 20001, ...

Resampling results:

Accuracy	Kappa
0.7225415	0.4396518

[Hide](#)

```
# Base Logistic Regression Test Prediction
lr_test_predict <- predict(model_lr, test_data)
confusionMatrix(lr_test_predict, test_data$Severity)
```

## Confusion Matrix and Statistics

```

    Reference
Prediction   0   1
      0 2060  764
      1  610 1565

    Accuracy : 0.7251
    95% CI  : (0.7125, 0.7375)
No Information Rate : 0.5341
P-Value [Acc > NIR] : < 2.2e-16

    Kappa : 0.4454

McNemar's Test P-Value : 3.666e-05

    Sensitivity : 0.7715
    Specificity  : 0.6720
    Pos Pred Value : 0.7295
    Neg Pred Value : 0.7195
    Prevalence   : 0.5341
    Detection Rate : 0.4121
    Detection Prevalence : 0.5649
    Balanced Accuracy : 0.7217

    'Positive' Class : 0

```

## Random Forest

[Hide](#)

```
# Load RF base model
model_rf <- readRDS("caret_rf_base.rds")
model_rf
```

Random Forest

```
20001 samples
 30 predictor
 2 classes: '0', '1'

No pre-processing
Resampling: Bootstrapped (25 reps)
Summary of sample sizes: 20001, 20001, 20001, 20001, 20001, ...
Resampling results across tuning parameters:
```

mtry	Accuracy	Kappa
2	0.8245709	0.6432299
16	0.8548752	0.7076566
30	0.8499837	0.6978825

Accuracy was used to select the optimal model using the largest value.  
The final value used for the model was mtry = 16.

[Hide](#)

```
# Base Random Forest Test Prediction
rf_test_predict <- predict(model_rf, test_data)
confusionMatrix(rf_test_predict, test_data$Severity)
```

### Confusion Matrix and Statistics

```
Reference
Prediction   0   1
  0 2348 367
  1 322 1962

Accuracy : 0.8622
95% CI : (0.8523, 0.8716)
No Information Rate : 0.5341
P-Value [Acc > NIR] : < 2e-16

Kappa : 0.7227

McNemar's Test P-Value : 0.09369

Sensitivity : 0.8794
Specificity : 0.8424
Pos Pred Value : 0.8648
Neg Pred Value : 0.8590
Prevalence : 0.5341
Detection Rate : 0.4697
Detection Prevalence : 0.5431
Balanced Accuracy : 0.8609

'Positive' Class : 0
```

## Feature Selection 3

### Feature Importance

[Hide](#)

```
importance_rf <- varImp(model_rf)
importance_rf
```

```
rf variable importance

only 20 most important variables shown (out of 30)
```

Overall
<dbl>

Weather_Timestamp	100.000000
Time	58.430005


Overall  
<dbl>

Zipcode	42.540973
Distance.mi.	23.323113
Pressure.in.	20.056902
State	19.847694
Street	18.687408
City	12.424100
Airport_Code	12.118603
Temperature.F.	11.965721

1-10 of 20 rows

Previous 1 2 Next

```
plot(varImp(model_rf))
```



```
# Convert importance object to dataframe
importance_df <- as.data.frame(importance_rf$importance)
importance_df
```

Overall  
<dbl>

Distance.mi.	23.3231127
Street	18.6874077
City	12.4241000
County	11.4904068
State	19.8476937
Zipcode	42.5409730
Timezone	6.1838766
Airport_Code	12.1186032
Weather_Timestamp	100.0000000
Temperature.F.	11.9657210

1-10 of 30 rows

Previous **1** 2 3 Next

[Hide](#)

```
# Sort features by importance (descending order)
sorted_features <- importance_df[order(-rowMeans(importance_df)), , drop = FALSE]

# Extract the top 20 features
top_features <- rownames(sorted_features)[1:20]
top_features
```

```
[1] "Weather_Timestamp" "Time"           "Zipcode"        "Distance.mi."
[5] "Pressure.in."      "State"          "Street"         "City"
[9] "Airport_Code"       "Temperature.F."   "County"        "Humidity..."
[13] "Wind_Speed.mph."   "Wind_Direction"  "Timezone"      "Weather_Condition"
[17] "Visibility.mi."    "Precipitation.in." "Sunrise_Sunset" "Junction"
```

[Hide](#)

```
### Extract New Train and Test Data based of Selected Features
selected_train_data <- train_data %>% select(all_of(top_features), Severity)
selected_test_data <- test_data %>% select(all_of(top_features), Severity)

# Check Any Missing Data in Training Data
anyNA(selected_train_data)
```

```
[1] FALSE
```

[Hide](#)

```
# Check Any Missing Data in Testing Data
anyNA(selected_test_data)
```

```
[1] FALSE
```

# Hyperparameter Tuning of Model

[Hide](#)

```
# Initialize Train Control  
trctrl <- trainControl(method = "repeatedcv",  
                         summaryFunction = defaultSummary, classProbs = FALSE,  
                         number = 5, repeats = 5, allowParallel = TRUE)
```

## Tune KNN

[Hide](#)

```
# Load KNN Tuned model  
model_knn_tune <- readRDS("caret_knn_tune.rds")  
print(model_knn_tune$bestTune)
```

k
<dbl>
14
27

1 row

[Hide](#)

```
model_knn_tune
```

## k-Nearest Neighbors

```
20001 samples  
20 predictor  
2 classes: '0', '1'
```

No pre-processing

Resampling: Cross-Validated (5 fold, repeated 5 times)

Summary of sample sizes: 16001, 16001, 16001, 16001, 16000, 16001, ...

Resampling results across tuning parameters:

k	Accuracy	Kappa
1	0.6998050	0.3953949
3	0.7209139	0.4363610
5	0.7263935	0.4466940
7	0.7296733	0.4526942
9	0.7301133	0.4530209
11	0.7307433	0.4540134
13	0.7313833	0.4550004
15	0.7319833	0.4560331
17	0.7323333	0.4564734
19	0.7324233	0.4564727
21	0.7335434	0.4585754
23	0.7342632	0.4599027
25	0.7340132	0.4592443
27	0.7347432	0.4606935
29	0.7338332	0.4587973
31	0.7342430	0.4594965
33	0.7339531	0.4588029
35	0.7338133	0.4584729
37	0.7331932	0.4571018
39	0.7333432	0.4573352
41	0.7336832	0.4579058
43	0.7336431	0.4577917
45	0.7340231	0.4584679
47	0.7343930	0.4591780
49	0.7336532	0.4575247
51	0.7329931	0.4561681
53	0.7326532	0.4554206
55	0.7321432	0.4542716
57	0.7320532	0.4540102
59	0.7314832	0.4527768
61	0.7317032	0.4531957
63	0.7312333	0.4522394
65	0.7307733	0.4512452
67	0.7312033	0.4521354
69	0.7313032	0.4523053
71	0.7314832	0.4526439
73	0.7311333	0.4519292
75	0.7309533	0.4515448
77	0.7314733	0.4525893
79	0.7309634	0.4514704
81	0.7307533	0.4510415
83	0.7304535	0.4504278
85	0.7305235	0.4505504

```
87  0.7300035  0.4494923
89  0.7301435  0.4497917
91  0.7296636  0.4487973
93  0.7303035  0.4500566
95  0.7305934  0.4506069
97  0.7299534  0.4493271
99  0.7295335  0.4484390
101 0.7294635  0.4482561
103 0.7295535  0.4484207
105 0.7292435  0.4478109
107 0.7295134  0.4483167
109 0.7295635  0.4484014
111 0.7298635  0.4490154
113 0.7294335  0.4480865
115 0.7293735  0.4479911
117 0.7288436  0.4469090
119 0.7286836  0.4465660
121 0.7285335  0.4462742
123 0.7286335  0.4464297
125 0.7284935  0.4461557
127 0.7280635  0.4452710
129 0.7281635  0.4454889
131 0.7277235  0.4445938
133 0.7272835  0.4437084
135 0.7270636  0.4432530
137 0.7269436  0.4429843
139 0.7270136  0.4431071
141 0.7271536  0.4433684
143 0.7269236  0.4428533
145 0.7266436  0.4422921
147 0.7266236  0.4422339
149 0.7260336  0.4410125
```

Accuracy was used to select the optimal model using the largest value.  
The final value used for the model was k = 27.

Hide

```
# Tuned KNN Test Prediction
knn_tune_test_predict <- predict(model_knn_tune, selected_test_data)
confusionMatrix(knn_tune_test_predict, selected_test_data$Severity)
```

## Confusion Matrix and Statistics

```

Reference
Prediction   0   1
      0 2236  817
      1  434 1512

Accuracy : 0.7497
95% CI  : (0.7375, 0.7617)
No Information Rate : 0.5341
P-Value [Acc > NIR] : < 2.2e-16

Kappa : 0.4918

McNemar's Test P-Value : < 2.2e-16

Sensitivity : 0.8375
Specificity  : 0.6492
Pos Pred Value : 0.7324
Neg Pred Value : 0.7770
Prevalence    : 0.5341
Detection Rate : 0.4473
Detection Prevalence : 0.6107
Balanced Accuracy : 0.7433

'Positive' Class : 0

```

## Tune Naive Bayes

Hide

```
# Load NB Tuned model
model_nb_tune <- readRDS("caret_nb_tune.rds")
print(model_nb_tune$bestTune)
```

fL	usekernel	adjust
<dbl>	<lgl>	<dbl>
4 0 TRUE 0.5		
1 row		

Hide

```
model_nb_tune
```

## Naive Bayes

```
20001 samples
 20 predictor
 2 classes: '0', '1'
```

No pre-processing

Resampling: Cross-Validated (5 fold, repeated 5 times)

Summary of sample sizes: 16001, 16001, 16001, 16001, 16000, 16001, ...

Resampling results across tuning parameters:

fL	usekernel	adjust	Accuracy	Kappa
0.0	FALSE	0.5	0.6656467	0.3133609
0.0	FALSE	1.0	0.6656467	0.3133609
0.0	FALSE	2.0	0.6656467	0.3133609
0.0	TRUE	0.5	0.7245439	0.4373821
0.0	TRUE	1.0	0.7189240	0.4238247
0.0	TRUE	2.0	0.7166742	0.4173014
0.5	FALSE	0.5	0.6656467	0.3133609
0.5	FALSE	1.0	0.6656467	0.3133609
0.5	FALSE	2.0	0.6656467	0.3133609
0.5	TRUE	0.5	0.7245439	0.4373821
0.5	TRUE	1.0	0.7189240	0.4238247
0.5	TRUE	2.0	0.7166742	0.4173014
1.0	FALSE	0.5	0.6656467	0.3133609
1.0	FALSE	1.0	0.6656467	0.3133609
1.0	FALSE	2.0	0.6656467	0.3133609
1.0	TRUE	0.5	0.7245439	0.4373821
1.0	TRUE	1.0	0.7189240	0.4238247
1.0	TRUE	2.0	0.7166742	0.4173014

Accuracy was used to select the optimal model using the largest value.

The final values used for the model were fL = 0, usekernel = TRUE and adjust = 0.5.

Hide

```
# Tuned Naive Bayes Test Prediction
nb_tune_test_predict <- predict(model_nb_tune, selected_test_data)
```

Warning: Numerical 0 probability for all classes with observation 492  
 Warning: Numerical 0 probability for all classes with observation 3008  
 Warning: Numerical 0 probability for all classes with observation 4564

Hide

```
confusionMatrix(nb_tune_test_predict, selected_test_data$Severity)
```

## Confusion Matrix and Statistics

```
Reference
Prediction   0   1
      0 2259  946
      1  411 1383

Accuracy : 0.7285
95% CI  : (0.716, 0.7408)
No Information Rate : 0.5341
P-Value [Acc > NIR] : < 2.2e-16

Kappa : 0.4464

McNemar's Test P-Value : < 2.2e-16

Sensitivity : 0.8461
Specificity  : 0.5938
Pos Pred Value : 0.7048
Neg Pred Value : 0.7709
Prevalence    : 0.5341
Detection Rate : 0.4519
Detection Prevalence : 0.6411
Balanced Accuracy : 0.7199

'Positive' Class : 0
```

## Tune Logistic Regression

[Hide](#)

```
# Load LR Tuned model
model_lr_tune <- readRDS("caret_lr_tune.rds")

# Predict with probability to find best threshold
train_probabilities <- predict(model_lr_tune, selected_train_data, type = "prob")

# set a dataframe to save results
lr_results <- data.frame()

# Set a probability cutoff
cutoffs <- c(0.3, 0.4, 0.5, 0.6, 0.7)

for (cutoff in cutoffs){
  train_predicted_class <- ifelse(train_probabilities[,2] > cutoff, 1, 0)

  # Convert to factor with correct levels
  train_predicted_class <- factor(train_predicted_class,
                                    levels = levels(selected_train_data$Severity))

  # Evaluate accuracy on training data
  cm <- confusionMatrix(train_predicted_class, selected_train_data$Severity)

  Accuracy <- cm$overall["Accuracy"]
  Kappa <- cm$overall["Kappa"]
  df_cm <- data.frame(cutoff, Accuracy, Kappa)
  lr_results <- rbind(lr_results, df_cm)
}

lr_results
```

	cutoff <dbl>	Accuracy <dbl>	Kappa <dbl>
Accuracy	0.3	0.6414679	0.3011257
Accuracy1	0.4	0.6875156	0.3786577
Accuracy2	0.5	0.7237638	0.4420431
Accuracy3	0.6	0.7363632	0.4604895
Accuracy4	0.7	0.7074646	0.3928161
5 rows			

[Hide](#)

```
# Tuned Logistic Regression Test Prediction

# Predict with probability to find best threshold
test_probabilities <- predict(model_lr_tune, selected_test_data, type = "prob")

# Set best probability cutoff
cutoff <- 0.6

test_predicted_class <- ifelse(test_probabilities[,2] > cutoff, 1, 0)

# Convert to factor with correct levels
test_predicted_class <- factor(test_predicted_class,
                                levels = levels(selected_test_data$Severity))

# Evaluate accuracy on training data
confusionMatrix(test_predicted_class, selected_test_data$Severity)
```

### Confusion Matrix and Statistics

		Reference	
Prediction	0	1	
0	2344	955	
1	326	1374	

Accuracy : 0.7437  
95% CI : (0.7314, 0.7558)  
No Information Rate : 0.5341  
P-Value [Acc > NIR] : < 2.2e-16

Kappa : 0.4761

McNemar's Test P-Value : < 2.2e-16

Sensitivity : 0.8779  
Specificity : 0.5900  
Pos Pred Value : 0.7105  
Neg Pred Value : 0.8082  
Prevalence : 0.5341  
Detection Rate : 0.4689  
Detection Prevalence : 0.6599  
Balanced Accuracy : 0.7339

'Positive' Class : 0

## Tune Random Forest

[Hide](#)

```
# Load RF base model
model_rf_tune <- readRDS("caret_rf_tune.rds")

# Print the best hyperparameters
print(model_rf_tune$finalModel)
```

Call:

```
randomForest(x = x, y = y, ntree = ..1, mtry = param$mtry, nodesize = ..2)
  Type of random forest: classification
  Number of trees: 500
  No. of variables tried at each split: 9
```

OOB estimate of error rate: 13.68%

Confusion matrix:

	0	1	class.error
0	9385	1297	0.1214192
1	1440	7879	0.1545230

[Hide](#)

model\_rf\_tune

Random Forest

20001 samples  
20 predictor  
2 classes: '0', '1'

No pre-processing

Resampling: Cross-Validated (5 fold, repeated 5 times)

Summary of sample sizes: 16000, 16001, 16002, 16000, 16001, 16000, ...

Resampling results across tuning parameters:

mtry	Accuracy	Kappa
1	0.8007602	0.5928636
2	0.8436876	0.6836665
3	0.8519171	0.7008603
4	0.8552171	0.7078208
5	0.8576869	0.7130234
6	0.8586569	0.7151143
7	0.8592369	0.7163680
8	0.8596869	0.7173261
9	0.8600068	0.7180478
10	0.8597370	0.7175212
11	0.8597670	0.7176259
12	0.8594970	0.7170772
13	0.8590370	0.7161796
14	0.8591570	0.7163863
15	0.8590971	0.7162769
16	0.8582971	0.7146896
17	0.8573872	0.7128362
18	0.8579672	0.7140331
19	0.8575073	0.7130929
20	0.8573172	0.7126973

Accuracy was used to select the optimal model using the largest value.

The final value used for the model was mtry = 9.

[Hide](#)

```
# Tuned Random Forest Test Prediction
rf_tune_test_predict <- predict(model_rf_tune, selected_test_data)
best_model_matrix <- confusionMatrix(rf_tune_test_predict, selected_test_data$Severity)
best_model_matrix
```

### Confusion Matrix and Statistics

		Reference
Prediction	0	1
0	2333	350
1	337	1979

Accuracy : 0.8626  
 95% CI : (0.8527, 0.872)  
 No Information Rate : 0.5341  
 P-Value [Acc > NIR] : <2e-16

Kappa : 0.7238


Mcnemar's Test P-Value : 0.6471

Sensitivity : 0.8738  
 Specificity : 0.8497  
 Pos Pred Value : 0.8695  
 Neg Pred Value : 0.8545  
 Prevalence : 0.5341  
 Detection Rate : 0.4667  
 Detection Prevalence : 0.5367  
 Balanced Accuracy : 0.8618

'Positive' Class : 0

[Hide](#)

```
# Plotting the confusion Matrix
model_conf_mat_data = as.data.frame(best_model_matrix$table)
model_conf_matrix_plot <- model_conf_mat_data %>%
  ggplot(aes(x = Prediction, y = Reference, fill = Freq)) +
  geom_tile() +
  geom_text(aes(label = Freq))
model_conf_matrix_plot
```



## Model Evaluation

[Hide](#)

```
# Plot ROC Curve
test_labels_binary <- as.numeric(selected_test_data$Severity) - 1
knn_model_binary <- as.numeric(knn_tune_test_predict) - 1
nb_model_binary <- as.numeric(nb_tune_test_predict) - 1
lr_model_binary <- as.numeric(test_predicted_class) - 1
rf_model_binary <- as.numeric(rf_tune_test_predict) - 1
```

[Hide](#)

```
# Compute the ROC Curve and AUC
roc_knn <- roc(test_labels_binary, knn_model_binary)
```

Setting levels: control = 0, case = 1  
 Setting direction: controls < cases

[Hide](#)

```
roc_nb <- roc(test_labels_binary, nb_model_binary)
```

Setting levels: control = 0, case = 1  
 Setting direction: controls < cases

[Hide](#)

```
roc_lr <- roc(test_labels_binary, lr_model_binary)
```

```
Setting levels: control = 0, case = 1  
Setting direction: controls < cases
```

[Hide](#)

```
roc_rf <- roc(test_labels_binary, rf_model_binary)
```

```
Setting levels: control = 0, case = 1  
Setting direction: controls < cases
```

[Hide](#)

```
# Plot all ROC curves together  
plot(roc_knn, col = "blue", lwd = 2, main = "ROC Curve Comparison")  
plot(roc_lr, col = "green", lwd = 2, add = TRUE)
```


[Hide](#)

```
plot(roc_nb, col = "purple", lwd = 2, add = TRUE)  
plot(roc_rf, col = "orange", lwd = 2, add = TRUE)
```

[Hide](#)

```
# Add legend  
legend("bottomright", legend = c(  
    paste("KNN (AUC =", round(auc(roc_knn), 2), ")"),  
    paste("LogReg (AUC =", round(auc(roc_lr), 2), ")"),  
    paste("Naive Bayes (AUC =", round(auc(roc_nb), 2), ")"),  
    paste("Random Forest (AUC =", round(auc(roc_rf), 2), ")")  
,  
    col = c("blue", "green", "purple", "orange"),  
    lwd = 2  
)
```

## ROC Curve Comparison


Hide

```
F1_knn <- F1_Score(test_labels_binary, knn_model_binary, positive = 0)
F1_nb <- F1_Score(test_labels_binary, nb_model_binary, positive = 0)
F1_lr <- F1_Score(test_labels_binary, lr_model_binary, positive = 0)
F1_rf <- F1_Score(test_labels_binary, rf_model_binary, positive = 0)

f1_scores <- data.frame(
  Model = c("KNN", "Naive Bayes", "Logistic Regression", "Random Forest"),
  F1 = c(F1_knn, F1_nb, F1_lr, F1_rf)
)
f1_scores
```

Model	F1
<chr>	<dbl>
KNN	0.7814084
Naive Bayes	0.7690213
Logistic Regression	0.7853912
Random Forest	0.8716608
4 rows	